



DevOps for Mere Mortals


How do smaller organizations run fast & smart?



Agenda

- Introductions
- Devops at Yahoo!
- Devops at SADA
- Devops for you

Allen Reese, Yahoo Inc.



YAHOO!

- JNI engineer at Yahoo!
- Moved Yahoo from
 - RHEL4 -> RHEL6 -> RHEL7
- Builds everyone else's packages.

Jack Weaver, SADA Systems



- Tech Lead at SADA
- Worked with Allen at Yahoo prior to SADA.
- Took SADA from “zero” to: devops, automated, tested, CI/CD, etc.

- All code, demos, documented details, are:
 - Deck: (Download to follow along, more readable).
 - <http://bit.ly/2qMobqx>
 - Github Repo:
 - <https://github.com/SwampUpMereMortals>
 - All code examples, as well as directions and setup notes.
 - Artifactory: <https://meremortal.jfrog.io/meremortal>
 - Travis: <https://travis-ci.org/SwampUpMereMortals>

“DevOps in the Small”: SADA Systems

- Company Size about 150 employees.
 - About 25 engineers.
 - Very successful company, traditionally consulting/IT services.
 - First production application “Atom”.
- Originally **zero** DevOps.
 - Engineers literally deploying to production from their personal laptop:
 - “What’s a runbook?”
 - No testing. No automation.
 - Outages went unnoticed.... Until the customer called us. No monitoring.
 - “Testing” was a human being clicking on the screen endlessly, day after day, week after week.
 - It’s amazing this person is still alive today. Medical science can’t explain it.
- Things are different now! Significant maturity in 12 months. *Serious* levels of automation.



“DevOps in the Small”: SADA Systems



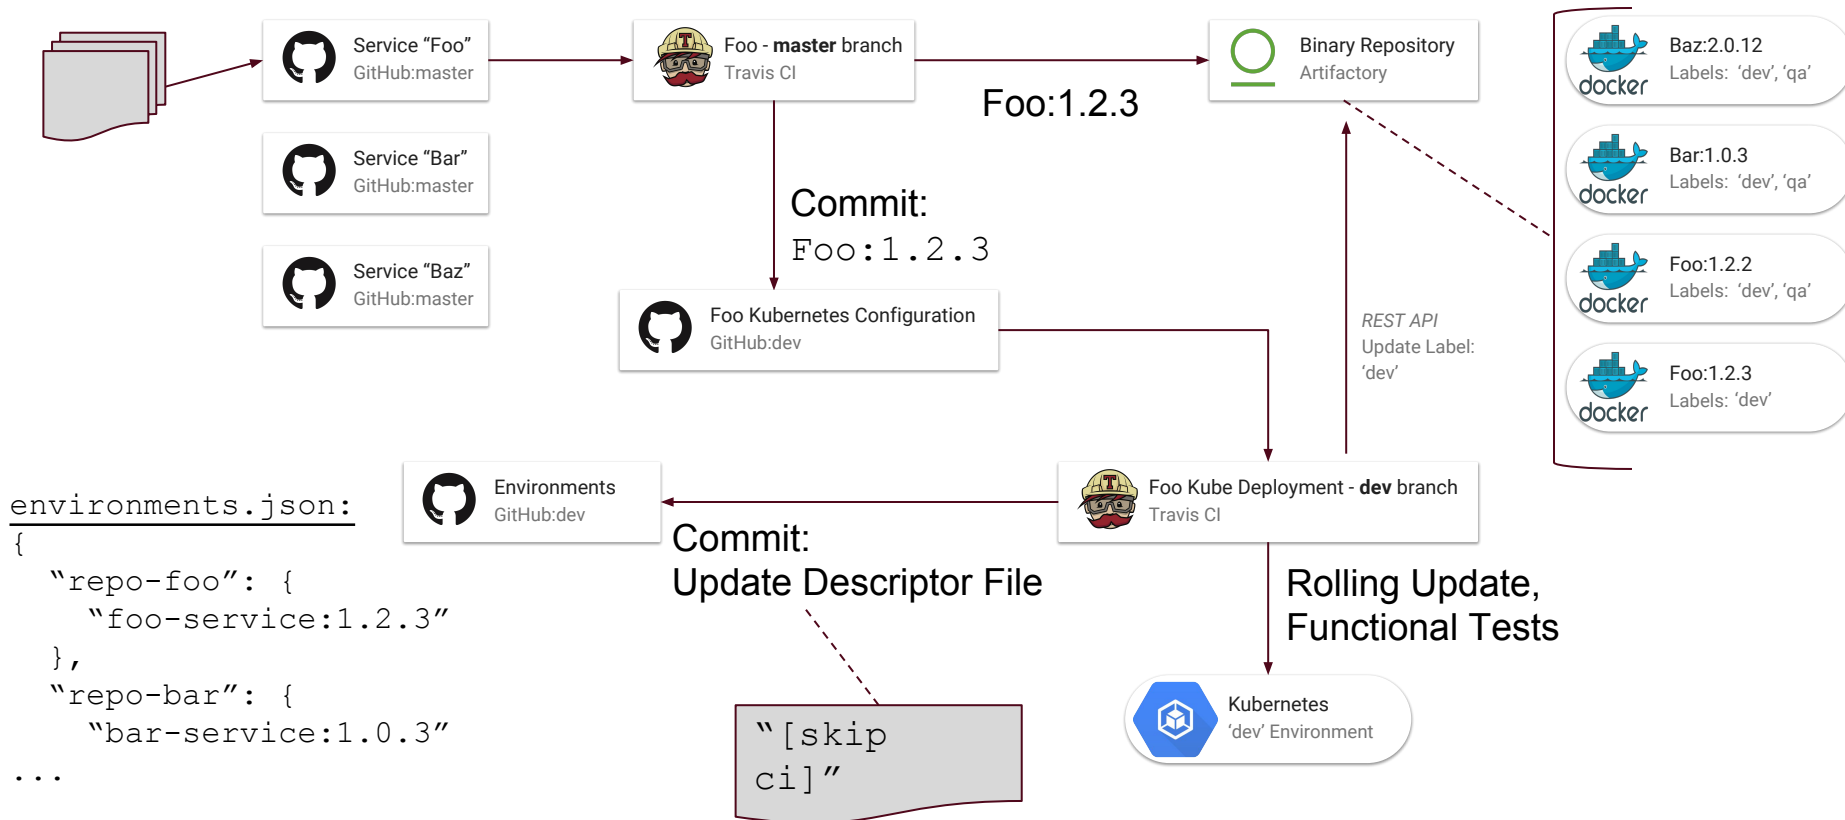
First principle: *software development is an economic activity.*

- Time spent on DevOps, Automation, CI/CD, etc, is an *investment* and thus must have an ROI.
- ***You therefore only invest in the ROIs that are greatest at your point in time.***

With a team size of 6 engineers, we started small and progressed outwards:

1. Setup Travis for all repos. Built, Unit Tested, Versioned & Published.
2. Infrastructure as Code from the start: code reviewed infrastructure changes, dedicated repos.
3. Started to want travis to deploy packages as well, setup CI commits.
4. Then wanted functional & integration tests, after deploy is done.
5. Finally, needed a way to coordinate deployment of multiple artifacts.

From Commit to Deployment



What does this do? Why?

- Deploy a single service, or an entire system (dev → qa, qa → prod).
- `environments.json` shows us exactly what is where.
- Single canonical version string, and everything is tagged.
- Different points in the pipeline correspond to different layers of testing: Unit Testing, Functional Testing, Integration Testing.
- Different points in the pipeline allow for interesting “articulation” points.
 - Need to know what was release into prod, from qa? Automate a CHANGELOG.

What the Big players do (Yahoo!)

- Large devtools teams to set everything up
 - 80+ people just to maintain build farms and tools to build it.
- Own the data centers
 - No clouds, no cost controls, no restrictions
 - Data is kept forever, until it causes a problem
- Control every aspect of what is run and deployed
 - Political and tool enforcement results in similar looking builds.
 - It's easier to build everything the way the other engineers understand
- Custom monitoring and CI tools built around Jenkins
 - Some of the choices around how to use Jenkins and sizing has resulted in poor scalability.
 - One box fits it all mentality.
- Internal SaaS platforms to push code through pipelines to production

What the Big players do (Yahoo!)

- Product teams make their own choices about builds.
 - On Jenkins every team has their own way of doing things.
 - With shared “build platforms” builds look similar.
 - Except when they don’t.
- Shared infrastructure
 - Shared outages.
 - Shared security holes
 - All source is effectively public
- Resource constraints to prevent monopolization of executors
 - Tried using docker, but cgroups broke the JDK.
 - JDK 8u131 has an option to behave better under docker
 - But you have to enable special flags to do it.
 - Which means modifying every build to know about the flags.
 - And the jvms that build tools launch

What can the little players do? (YOU!)



- Rely on cloud providers for infrastructure
 - No hardware to buy/maintain, no datacenters to put the hardware in
 - Fine grained cost controls.
- Use SaaS solutions to meet your needs.
 - Bintray, Artifactory SaaS, Travis CI/Circle/Drone/Semaphore/etc, GCP/Google Container Engine, StackDriver, Codacy
 - Sometimes buggy, or limited in ways that are hard to work around
 - Their downtime means you are down. It's a risk.

What can the little players do? (YOU!)



- Limit the amount of engineering time required to push code to production
 - Leverage travis or circle.
 - Can require lots of glue code. Everything is an ROI, measure it before you put it in a sprint.
 - Avoid unknown outages and large manual QA effort.
 - Smaller teams necessitate the need for automation.
- Your Advantage: new development, no legacy debt. Can design software and architecture using modern development practices.
 - Lower ramp up times for new engineers as they may already know these.

What can the little players do? (YOU!)



- Understand what tools you are using and why
 - If you don't know why you are using a tool, or how others use it, you might be using the wrong tool.
- Understand how the parts fit together.
 - Sometimes builds get so complicated, you can only guess if it works by committing the code.
 - This prevents debugging, and frustrates your developers.
 - Complicated build systems break in interesting ways. The more the people using the build can read and understand about how things fit together, the faster code goes to production.
- Understand how the solution fits your situation
 - For a devops system to work correctly, you need to understand how the system influences your end goal of increased reliability and decreased downtime.

Why change what you are doing?

- The root cause of a large number of incidents was found to be human error.
- People accidentally mis-key or skip steps and tests in the runbooks.
- Each engineer deploys code differently.
- Allows traceback of what code is running in production.
 - A site was running 39 different jars in a 40 node cluster for a single service.
 - 38 of those jars were built on developer's boxes from code that was never committed

LIVE DEMO

Let's hope this works.



Artifactory SaaS on GCP


- Setup Artifactory SaaS in GCP
 - We'll be deploying our software to artifactory.
 - <http://meremortal.jfrog.io/meremortal>
- Free 30 day trial
- *Protip*: use same GCP region as GKE.
- Gradle/maven/docker work seamlessly with artifactory SaaS or on-prem.
 - <https://github.com/JFrogDev/project-examples>



JFrog Artifactory

Travis CI Setup

- Setup TravisCI to deploy to Artifactory
 - We'll use Travis as our CI provider.
- Installs common-ci (debian package)
- Configure for Service Accounts (GCP) & github GH_TOKEN (for auth/commits back)
- Room for optimization (caching).



Multiple Repos
Travis CI

<https://travis-ci.org/SwampUpMereMortals>

NOTE: You can easily use Circle, Semaphore, Drone, etc. We chose Travis randomly.

Kubernetes on GCP/GKE

- Setup Kubernetes within GCP
 - We'll be deploying to the cloud as well.
- Rules of the road:
 - We separate out application code from deployment code. Those are two different concerns.
 - Application code repo & Kubernetes Pod/Deployment repo
 - Immutable Server Pattern
 - <https://martinfowler.com/bliki/ImmutableServer.html>
 - CI/CD does everything for us. Minimize the opportunity for fat fingers and human mistakes.



The Top 4 Takeaways for Mere Mortals

What can You learn from this?



Automate: Our most active contributor



Atom CI
atom-headless-ci

<http://pastebin.com/raw/NADFxGx>

Unfollow

Block or report user

@sadasystems



Overview

Repositories 0

Stars


Followers 2

Following 0

Popular repositories

atom-headless-ci doesn't have any repositories you can view.

674 contributions in the last year



Contribution activity

Jump to ▾

2017

May 2017

2016



Created 121 commits in 8 repositories



- Immutable Server Pattern.
- Cattle, not pets.
- No snowflakes in your system.
- Don't let data and state dictate your systems.
 - Stateful sets or pet sets in Kubernetes.

Keep everything simple and similar

- The simpler the setup is, the easier it is to describe and modify
- Keep things similar without copying too much boilerplate
- Share where it makes sense, the more you share the more you are forced to keep things simple.
- Simple things are easy to debug by sight.

- Only invest your time in the biggest payoffs for your organization at the time.
 - Don't build the world. Start **small**. Receive immediate returns on that time.
 - For example, in progression:
 - Infrastructure as Code (review it). Know your hardware & configurations.
 - CI Builds, Unit Test everything at build time.
 - Start versioning, packaging & publishing from CI builds.
 - Start having CI conduct deployments.
 - Start running Functional/Integration tests after deployment.
- Use Monitoring & Alerting Systems (sometimes you get this free).
- Use automation where it makes sense for you, this is not one size fits all.



**Please submit
feedback forms!**

Questions?